

# Similarity: Regression

Valari Graham

2022-10-08

## Regression

This notebook performs regression using a linear model, kNN, and decision trees. # Data Set Link to data set in UCI Machine Learning Repository “<https://archive.ics.uci.edu/ml/datasets/Bike+Sharing+Dataset>”

Load in data set and remove column ‘date’

```
library(tree)
library(MASS)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following object is masked from 'package:MASS':
##
##      select

## The following objects are masked from 'package:stats':
##
##      filter, lag

## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union
```

```
library(ggplot2)
bike1 <- read.csv("hour.csv")
bike <- select(bike1, -c(dteday, season, weathersit))
```

## Linear Regression Model

Split into train and test data

```
set.seed(1234)
i <- sample(nrow(bike), 0.8*nrow(bike), replace = FALSE)
train <- bike[i,]
test <- bike[-i,]
```

Brief Statistical Exploration of the Data

```
str(train)
```

```
## 'data.frame': 13903 obs. of 14 variables:
## $ instant : int 7452 8016 7162 8086 9196 623 15241 10885 934 12688 ...
## $ yr : int 0 0 0 0 1 0 1 1 0 1 ...
## $ mnth : int 11 12 10 12 1 1 10 4 2 6 ...
## $ hr : int 2 15 0 13 1 4 5 16 12 20 ...
## $ holiday : int 0 0 0 0 0 0 0 0 0 0 ...
## $ weekday : int 6 1 1 4 2 6 2 2 5 0 ...
## $ workingday: int 0 1 1 1 1 0 1 1 1 0 ...
## $ temp : num 0.24 0.46 0.26 0.3 0.32 0.16 0.56 0.62 0.22 0.62 ...
## $ atemp : num 0.258 0.455 0.303 0.273 0.303 ...
## $ hum : num 0.65 0.72 0.87 0.49 0.93 0.69 0.83 0.21 0.47 0.57 ...
## $ windspeed : num 0.0896 0.0896 0 0.3582 0.2537 ...
## $ casual : int 7 16 3 9 2 1 1 145 7 101 ...
## $ registered: int 39 132 20 115 7 2 42 340 64 201 ...
## $ cnt : int 46 148 23 124 9 3 43 485 71 302 ...
```

```
summary(train)
```

```
## instant yr mnth hr
## Min. : 1 Min. :0.0000 Min. : 1.000 Min. : 0.00
## 1st Qu.: 4386 1st Qu.:0.0000 1st Qu.: 4.000 1st Qu.: 6.00
## Median : 8768 Median :1.0000 Median : 7.000 Median :12.00
## Mean : 8730 Mean :0.5064 Mean : 6.549 Mean :11.55
## 3rd Qu.:13054 3rd Qu.:1.0000 3rd Qu.:10.000 3rd Qu.:18.00
## Max. :17379 Max. :1.0000 Max. :12.000 Max. :23.00
## holiday weekday workingday temp
## Min. :0.00000 Min. :0.000 Min. :0.00000 Min. :0.0200
## 1st Qu.:0.00000 1st Qu.:1.000 1st Qu.:0.00000 1st Qu.:0.3400
## Median :0.00000 Median :3.000 Median :1.00000 Median :0.5000
## Mean :0.02877 Mean :3.011 Mean :0.6821 Mean :0.4972
## 3rd Qu.:0.00000 3rd Qu.:5.000 3rd Qu.:1.00000 3rd Qu.:0.6600
## Max. :1.00000 Max. :6.000 Max. :1.00000 Max. :1.0000
## atemp hum windspeed casual
## Min. :0.0000 Min. :0.0000 Min. :0.0000 Min. : 0.00
## 1st Qu.:0.3333 1st Qu.:0.4800 1st Qu.:0.1045 1st Qu.: 4.00
## Median :0.4848 Median :0.6300 Median :0.1940 Median : 17.00
## Mean :0.4760 Mean :0.6265 Mean :0.1896 Mean : 35.89
## 3rd Qu.:0.6212 3rd Qu.:0.7800 3rd Qu.:0.2537 3rd Qu.: 49.00
## Max. :0.9848 Max. :1.0000 Max. :0.8507 Max. :367.00
## registered cnt
## Min. : 0.0 Min. : 1.0
## 1st Qu.: 35.0 1st Qu.: 41.0
## Median :117.0 Median :144.0
## Mean :154.9 Mean :190.8
## 3rd Qu.:221.0 3rd Qu.:282.0
## Max. :886.0 Max. :977.0
```

```
head(train)
```

```
## instant yr mnth hr holiday weekday workingday temp atemp hum windspeed
```


```
## 7452    7452  0    11  2      0      6      0 0.24 0.2576 0.65    0.0896
## 8016    8016  0    12 15      0      1      1 0.46 0.4545 0.72    0.0896
## 7162    7162  0    10  0      0      1      1 0.26 0.3030 0.87    0.0000
## 8086    8086  0    12 13      0      4      1 0.30 0.2727 0.49    0.3582
## 9196    9196  1     1  1      0      2      1 0.32 0.3030 0.93    0.2537
## 623      623  0     1  4      0      6      0 0.16 0.1818 0.69    0.1045
##          casual registered cnt
## 7452         7          39 46
## 8016        16         132 148
## 7162         3          20 23
## 8086         9         115 124
## 9196         2           7  9
## 623         1           2  3
```

```
sum(is.na(train))
```

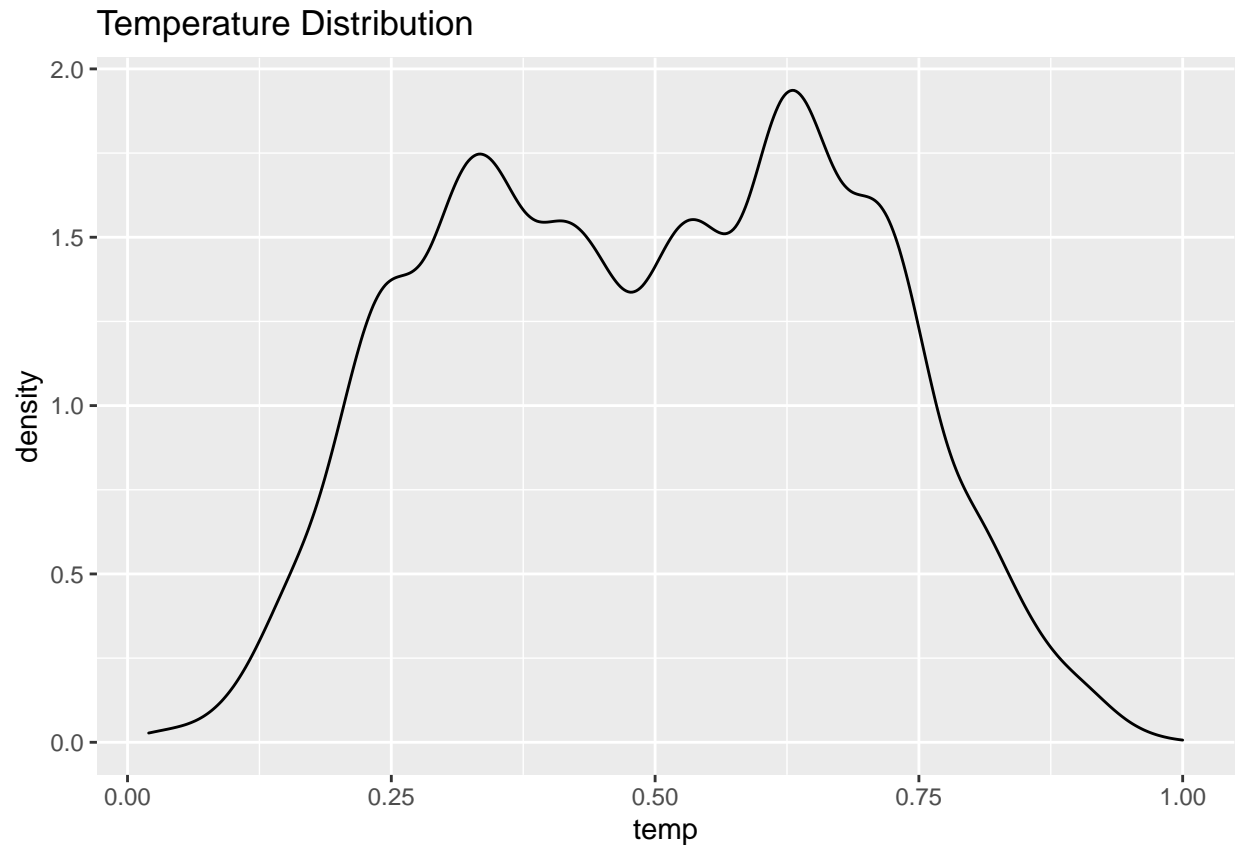
```
## [1] 0
```

Graphs on the data

```
ggplot(train, aes(x = cnt)) + geom_density() + ggtitle("Count of Bikes Rented Density Plot")
```



```
ggplot(train, aes(x = temp)) + geom_density() + ggtitle("Temperature Distribution")
```



Create base linear model

```
lm1 <- lm(cnt~., data=train)
summary(lm1)
```

```
##
## Call:
## lm(formula = cnt ~ ., data = train)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-4.641e-11	-1.300e-14	7.000e-15	2.800e-14	2.723e-11

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-2.706e-13	3.700e-14	-7.314e+00	2.73e-13 ***
instant	3.722e-17	2.537e-17	1.467e+00	0.142424
yr	-8.197e-13	2.222e-13	-3.689e+00	0.000226 ***
mnth	-9.831e-14	1.851e-14	-5.310e+00	1.11e-07 ***
hr	1.481e-14	8.603e-16	1.721e+01	< 2e-16 ***
holiday	8.127e-14	3.316e-14	2.451e+00	0.014271 *
weekday	1.915e-16	2.676e-15	7.200e-02	0.942948
workingday	-4.018e-13	1.353e-14	-2.969e+01	< 2e-16 ***
temp	1.809e-12	1.809e-13	1.000e+01	< 2e-16 ***
atemp	9.408e-13	2.033e-13	4.628e+00	3.72e-06 ***
hum	-6.313e-13	3.205e-14	-1.970e+01	< 2e-16 ***

```
## windspeed    7.446e-14  4.737e-14  1.572e+00 0.115982
## casual       1.000e+00  1.586e-16  6.304e+15 < 2e-16 ***
## registered   1.000e+00  4.646e-17  2.152e+16 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.281e-13 on 13889 degrees of freedom
## Multiple R-squared:      1, Adjusted R-squared:      1
## F-statistic: 8.99e+31 on 13 and 13889 DF, p-value: < 2.2e-16
```

```
pred <- predict(lm1, newdata=test)
cor_lm <- cor(pred, test$cnt)
mse_lm <- mean((pred - test$cnt)^2)
rmse_lm <- sqrt(mean((pred-test$cnt)^2))
print(paste("cor=", cor_lm))
```

```
## [1] "cor= 1"
```

```
print(paste("mse=", mse_lm))
```

```
## [1] "mse= 4.07687806854585e-25"
```

The R squared value measures the relationship between your linear model and target variable. The results show the linear model has a perfect fit with an extremely low mean squared error. The independent variables overall have very low p-values as well.

## kNN Regression

Scale data and run regression

```
train_scaled <- train[,1:14]
means <- sapply(train_scaled, mean)
stdvs <- sapply(train_scaled, sd)
train_scaled <- scale(train_scaled, center=means, scale=stdvs)
test_scaled <- scale(test[,1:14], center=means, scale=stdvs)
```

Run kNN Regression

```
library(caret)
```

```
## Loading required package: lattice
```

```
fit <- knnreg(train_scaled, train$cnt, k=3)
pred2 <- predict(fit, test_scaled)
cor_knn2 <- cor(pred2, test$cnt)
mse_knn2 <- mean((pred2 - test$cnt)^2)
print(paste("cor=", cor_knn2))
```

```
## [1] "cor= 0.989948580342323"
```

```
print(paste("mse=", mse_knn2))
```

```
## [1] "mse= 642.832951748498"
```

```
print(paste("rmse=", sqrt(mse_knn2)))
```

```
## [1] "rmse= 25.3541505822715"
```

Find the best k value for the model

```
cor_k <- rep(0, 20)
mse_k <- rep(0, 20)
i <- 1
for (k in seq(1, 39, 2)){
  fit_k <- knnreg(train_scaled, train$cnt, k=k)
  pred_k <- predict(fit_k, test_scaled)
  cor_k[i] <- cor(pred_k, test$cnt)
  mse_k[i] <- mean((pred_k - test$cnt)^2)
  print(paste("k=", k, cor_k[i], mse_k[i]))
  i <- i + 1
}
```


```
## [1] "k= 1 0.98306775875943 1086.10241657077"
## [1] "k= 3 0.989948580342323 642.832951748498"
## [1] "k= 5 0.991008627276036 577.514972509909"
## [1] "k= 7 0.991078196656088 578.53195105358"
## [1] "k= 9 0.990837665291767 598.167699924704"
## [1] "k= 11 0.990687045758844 610.420692204237"
## [1] "k= 13 0.990536133453734 622.098251345261"
## [1] "k= 15 0.990317061453034 640.608916618719"
## [1] "k= 17 0.990101790540697 657.122606021795"
## [1] "k= 19 0.989826628747146 677.897159308069"
## [1] "k= 21 0.98953047445573 701.405957997914"
## [1] "k= 23 0.989391229072734 714.556336770811"
## [1] "k= 25 0.989123702547015 736.736215366912"
## [1] "k= 27 0.988916907999926 756.562705016046"
## [1] "k= 29 0.988689256183867 775.086251720459"
## [1] "k= 31 0.988474479915493 791.995322131275"
## [1] "k= 33 0.988231318992835 811.468623235692"
## [1] "k= 35 0.987965197510416 831.87150405248"
## [1] "k= 37 0.987713085061437 852.550849626071"
## [1] "k= 39 0.987527413132629 868.009269547041"
```

```
plot(1:20, cor_k, lwd=2, col='red', ylab="", yaxt='n')
par(new=TRUE)
plot(1:20, mse_k, lwd=2, col='blue', labels=FALSE, ylab="", yaxt='n')
```

```
## Warning in plot.window(...): "labels" is not a graphical parameter
```

```
## Warning in plot.xy(xy, type, ...): "labels" is not a graphical parameter
```

```
## Warning in box(...): "labels" is not a graphical parameter
## Warning in title(...): "labels" is not a graphical parameter
```



Find when correlation is the highest and mse is the lowest.

```
which.min(mse_k)
```

```
## [1] 3
```

```
which.max(cor_k)
```

```
## [1] 4
```

kNN with  $k = 4$

```
fit <- knnreg(train_scaled, train$cnt, k=4)
pred3 <- predict(fit, test_scaled)
cor_knn3 <- cor(pred3, test$cnt)
mse_knn3 <- mean((pred3 - test$cnt)^2)
print(paste("cor=", cor_knn3))
```

```
## [1] "cor= 0.990648451008266"
```

```
print(paste("mse=", mse_knn3))
```

```
## [1] "mse= 599.040376150748"
```

The results for kNN regression has a lower correlation than the linear model. As well as, it has a much higher mean squared error than the previous linear model. The linear model out performed the bike rental data set. `## Decisions Trees Using Regression`

```
tree1 <- tree(cnt~., data=train)
summary(tree1)
```

```
##
## Regression tree:
## tree(formula = cnt ~ ., data = train)
## Variables actually used in tree construction:
## [1] "registered" "casual"
## Number of terminal nodes: 9
## Residual mean deviance: 1555 = 21600000 / 13890
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -149.20  -18.48   -4.03    0.00   17.37   230.80
```

```
pred <- predict(tree1, newdata=test)
print(paste('correlation:', cor(pred, test$cnt)))
```

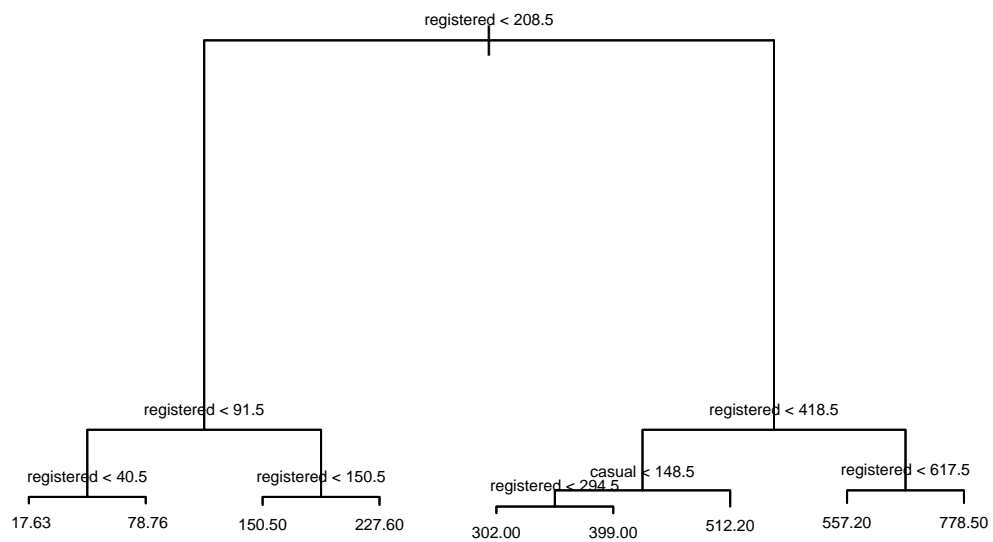
```
## [1] "correlation: 0.97454610774393"
```

```
rmse_tree <- sqrt(mean((pred-test$cnt)^2))
print(paste('rmse:', rmse_tree))
```


```
## [1] "rmse: 40.0014591375194"
```

```
plot(tree1)
text(tree1, cex=0.5, pretty=0)
```






```
cv_tree <- cv.tree(tree1)
plot(cv_tree$size, cv_tree$dev, type='b')
```



The results show that that the best size of the tree is 9. Proceed with pruning.

```
tree_pruned <- prune.tree(tree1, best=9)
plot(tree_pruned)
text(tree_pruned, pretty=0)
```



### Pruning Results

```

pred_pruned <- predict(tree_pruned, newdata=test)
cor_pruned <- cor(pred_pruned, test$cnt)
rmse_pruned <- rmse_pruned <- sqrt(mean((pred_pruned-test$cnt)^2))
print(paste("correlation:", cor_pruned))

```

```
## [1] "correlation: 0.97454610774393"
```

```
print(paste("rmse:", rmse_pruned))
```

```
## [1] "rmse: 40.0014591375194"
```

The results for the decision tree after pruning have a lower correlation than the linear model and kNN regression. The mean squared error for the decision tree is lower than the result of mse for the kNN regression.

Overall, the linear model out-performs both the kNN regression and decision tree. This results goes to show that sometimes more complex models are not always the right solution for data analysis and exploration. When more basic models perform well, staticians use less resources to perform analysis.